

PANNAGARAM P R

Shimoga, Karnataka, India. — pannagarampr@gmail.com — +91 6366206533
linkedin.com/in/pannagarampr — github.com/Pannagaram

Education

Manipal School of Information Sciences, MAHE <i>Master of Engineering in Big Data Analytics (CGPA: 8.44)</i>	Manipal, India July 2025 – Present
R N Shetty Institute of Technology <i>Bachelor of Engineering in Artificial Intelligence and Machine Learning (CGPA: 8.49)</i>	Bengaluru, India Dec 2021 – May 2025
Pace UM Pre University College <i>Pre University Course (PCMB) (Aggregate: 100)</i>	Shimoga, India June 2019 – May 2021
Saandeepani English High School <i>General (Aggregate: 96.64)</i>	Shimoga, India June 2018 – April 2019

Skills Summary

Programming Languages: Python, R

Libraries / Frameworks: NumPy, Pandas, Scrapy, BeautifulSoup, PySpark, Streamlit

Machine Learning / Deep Learning: Transformers, TensorFlow, Scikit-learn, Hugging Face, Pinecone, LangChain

Tools / Platforms: Hadoop, Sqoop, Spark, Hive, Airflow, Docker, Jenkins, PowerBI

Databases: MySQL

Coursework: Architecture of Big Data, Machine Learning, Data Structures and Algorithms, Operating Systems, Computer Networks, Computer Organization and Architecture

Projects

MSIS Placement Assistant – Agentic RAG Chatbot <i>Agentic RAG Based Chatbot</i>	April 2026 – May 2026
- Designed an end-to-end placement intelligence workflow processing company posters and JDs via OCR, text cleaning, chunking, and metadata enrichment, forming a structured knowledge base for downstream retrieval. - Implemented an Agentic RAG system combining LLM reasoning, tool orchestration, and multi-step planning to autonomously retrieve, analyze, and respond to placement queries using contextual document intelligence. - Engineered RAG retrieval pipelines using Sentence Transformers and ChromaDB with role-based filtering, integrated into a Streamlit interface.	
Wikipedia Trend Mining ETL Pipeline <i>Airflow Orchestrated ETL Pipeline</i>	Nov 2025 – Dec 2025
- Built an end-to-end data engineering pipeline to analyze trending Wikipedia topics by ingesting large-scale pageview data from the Wikimedia Pageview API and orchestrating workflows using Apache Airflow and PySpark. - Designed a Medallion Architecture (Bronze–Silver–Gold) data pipeline on a MinIO data lake, enabling structured data processing through staged layers for raw ingestion, cleaned transformations, and analytical datasets. - Containerized the pipeline using Docker and developed interactive Power BI dashboards to visualize trending topics, traffic patterns, and insights derived from Wikipedia pageview data.	
Wind Energy Forecasting Pipeline on Azure using MLflow <i>Machine Learning Forecasting System</i>	Jan 2026 - Feb 2026
- Developed a Neural Network based wind power forecasting model using TensorFlow Keras for energy prediction based on weather features. - Managed the ML lifecycle using MLflow on Azure Databricks, including experiment tracking, model versioning, and deployment. - Implemented CI/CD pipelines using GitHub Actions for automated model training and deployment.	

Certifications and Technical Achievements

Snowflake Badges — Data Warehouse, Data Applications, Data Engineering (Sept 2025)
Generative AI with Large Language Models — Coursera (Nov 2025)
AWS Academy Graduate — Cloud Foundations (Dec 2025)